

Probabilities
Expectation
 $\mathbb{E}[X] = \int_{\Omega} xf(x)dx = \int_{\omega} xP[X=x]dx$
 $\mathbb{E}_{Y|X}[Y] = \mathbb{E}_Y[Y|X]$
 $\mathbb{E}_{X,Y}[f(X,Y)] = \mathbb{E}_X\mathbb{E}_{Y|X}[f(X,Y)|X]$
 $\mathbb{E}_{Y|X}[f(X,Y)|X] = \int_{\mathbb{R}} f(X,y)p_{Y|X}(y)dy$
Variance & Covariance
 $\text{Var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$
 $\text{Var}[aX \pm bY] = a^2\text{Var}[X] + b^2\text{Var}[Y] \pm 2ab\text{Cov}[X,Y]$ $XY \text{ iid}$
 $\text{Cov}(X,Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$
Conditional Probabilities
 $P[X|Y] = \frac{P[X,Y]}{P[Y]}, P[\bar{X}|Y] = 1 - P[X|Y]$
Distributions
 $\mathcal{N}(x|\mu, \sigma^2) = 1/(\sqrt{2\pi\sigma^2})e^{-(x-\mu)^2/(2\sigma^2)}$
 $\mathcal{N}(x|\mu, \Sigma) = \frac{1}{(2\pi)^{2D}/|\Sigma|^{1/2}} \text{mathrme}^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$
 $\text{Exp}(x|\lambda) = \lambda e^{-\lambda x}, \text{Ber}(x|\theta) = \theta^x(1-\theta)^{(1-x)}$
Sigmoid: $\sigma(x) = 1/(1 + \exp(-x))$
Chebyshev & Consistency
 $P(|X - \mathbb{E}[X]| \geq \epsilon) \leq \frac{\text{Var}(X)}{\epsilon^2}$
 $\lim n \rightarrow \infty P(|\hat{\mu} - \mu| > \epsilon) = 0$
Cramer Rao lower bound
 $\text{Var}[\hat{\theta}] \geq \mathcal{I}_n(\theta)$
 $\mathcal{I}_n(\theta) = -\mathbb{E}[\frac{\partial^2 \log[\mathcal{L}_n|\theta]}{\partial \theta^2}]$ $\hat{\theta}$ unbiased
Efficiency of $\hat{\theta}$: $e(\theta_n) = \frac{1}{\text{Var}[\hat{\theta}_n|\mathcal{I}_n(\theta)]}$
 $e(\theta_n) = 1$ (efficient)
 $\lim_{n \rightarrow \infty} e(\theta_n) = 1$ (asympt. efficient)
Matrix Derivations
 $\frac{\partial \mathbf{a}^T \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}$ $\frac{\partial \mathbf{a}^T \mathbf{X} \mathbf{b}}{\partial \mathbf{X}} = \mathbf{a} \mathbf{b}^T$ $\frac{\partial \mathbf{a}^T \mathbf{X}^T \mathbf{b}}{\partial \mathbf{X}} = \mathbf{b} \mathbf{a}^T$
 $\frac{\partial \mathbf{x}^T \mathbf{A} \mathbf{x}}{\partial \mathbf{x}} = (\mathbf{A} + \mathbf{A}^T) \mathbf{x}$
 $\frac{\partial}{\partial \mathbf{x}} \mathbf{f}^T \mathbf{g} = \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \mathbf{g} + \mathbf{g}^T \left(\frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right)^T$
 $\mathbf{X}^T \mathbf{X}$: only invertible if none of the Eigenvalue is 0. Inversion instable if ratio from \mathbf{X} 's smallest EV to the largest is big.
Optimization
Gradient Descent
 $\theta^{\text{new}} \leftarrow \theta^{\text{old}} - \eta \nabla_{\theta} \mathcal{L}$
Convergence isn't guaranteed.
Less zigzag by adding momentum:
 $\theta^{(l+1)} \leftarrow \theta^{(l)} - \eta \nabla_{\theta} \mathcal{L} + \mu(\theta^{(l)} - \theta^{(l-1)})$
Newton's Method
Use 2nd order derivation. (Hessian)
 $\theta^{\text{new}} \leftarrow \theta^{\text{old}} - \eta (\nabla_{\theta} \mathcal{L} / \nabla_{\theta}^2 \mathcal{L})$
 $H = \nabla_{\theta}^2 \mathcal{L}$ has to be p.d (convex func).

Risks and Losses
Expected Risk
Conditional Expected Risk
 $R(f, X) = \int_{\mathbb{R}} \mathcal{L}(Y, f(X))P(Y|X)dY$
Total Expected Risk $R(f) = \mathbb{E}_X[R(f, X)] = \int_{\mathcal{X}} R(f, X)P(X)dX = \int_{\mathcal{X}} \int_{\mathbb{R}} \mathcal{L}(Y, f(X))P(X, Y)dXdY$
Empirical Risk
 $Z^{\text{train}} = (X_1, Y_1), \dots, (X_n, Y_n)$
 $Z^{\text{test}} = (X_{n+1}, Y_{n+1}), \dots, (X_{n+m}, Y_{n+m})$
Empirical Risk Minimizer \hat{f} s.t.
 $\hat{f} \in \arg \min_{f \in \mathcal{C}} \hat{R}(\hat{f}, Z^{\text{train}})$
Training error:
 $\hat{R}(\hat{f}, Z^{\text{train}}) = \frac{1}{n} \sum_{i=1}^n Q(Y_i, \hat{f}(X_i))$
Test error:
 $\hat{R}(\hat{f}, Z^{\text{test}}) = \frac{1}{m} \sum_{i=n+1}^{n+m} Q(Y_i, \hat{f}(X_i))$
 $\hat{R}(\hat{f}, Z^{\text{test}}) \neq \mathbb{E}_X[R(f, X)]$
Linear Regression
Data: $Z = (x_i, y_i) \in \mathbb{R}^3 \times \mathbb{R} : 1 \leq i \leq n$
 X are iids and Y depends on X .
Model: $\mathbf{Y} = \beta_0 + \sum_{j=1}^d \mathbf{X}_j \beta_j$ $\mathbf{Y} \subset \mathbb{R}$
Introduce $X_0 = 1$ and rewrite
 $\mathbf{Y} = \mathbf{X}^T \boldsymbol{\beta}$ $\mathbf{X} \in \mathbb{R}^{(d+1) \times n}, \boldsymbol{\beta} \in \mathbb{R}^{d+1}$
additive Gaussian noise $\epsilon \sim \mathcal{N}(0, \sigma^2)$
 $\hat{\mathbf{y}} = \mathbf{X} \hat{\boldsymbol{\beta}} + \epsilon$
 $\hat{\boldsymbol{\beta}} \sim \mathcal{N}(\boldsymbol{\beta}, (\mathbf{X}^T \mathbf{X})^{-1} \sigma^2)$ and
 $p(Y|X, \boldsymbol{\beta}, \sigma) \sim \mathcal{N}(Y|X^T \boldsymbol{\beta}, \sigma^2)$
A Regression has Optimum:
 $f^*(x) = \mathbb{E}_Y[Y|X=x]$
Linear Regression
Setting: Minimize RSS.
 $\mathcal{L} = \text{RSS}(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - x_i^T \boldsymbol{\beta})^2 = (\mathbf{y} - \mathbf{X} \boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X} \boldsymbol{\beta})$
 $\mathcal{L} = \sum_{i=1}^n (y_i - x_i^T \boldsymbol{\beta})^2 + \lambda \sum_{j=1}^d \beta_j^2 = (\mathbf{y} - \mathbf{X} \boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X} \boldsymbol{\beta}) + \lambda \boldsymbol{\beta}^T \boldsymbol{\beta}$
Solution: differentiate w.r.t $\boldsymbol{\beta}$
 $\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
Is an orth. projection with lowest variance of all unbiased estimates.
Prediction: $\hat{\mathbf{y}} = \mathbf{X} \hat{\boldsymbol{\beta}} = \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
Ridge Regression (L2 penalty)
Setting: Penalize the $\boldsymbol{\beta}$ s
 $\mathcal{L} = \sum_{i=1}^n (y_i - x_i^T \boldsymbol{\beta})^2 + \lambda \sum_{j=1}^d \beta_j^2 = (\mathbf{y} - \mathbf{X} \boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X} \boldsymbol{\beta}) + \lambda \boldsymbol{\beta}^T \boldsymbol{\beta}$
Solution: differentiate w.r.t $\boldsymbol{\beta}$
 $\hat{\boldsymbol{\beta}}^{\text{ridge}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$
Lasso (L1 penalty)
Setting: seek for a sparse solution
 $\mathcal{L} = \sum_{i=1}^n (y_i - x_i^T \boldsymbol{\beta})^2 + \lambda \sum_{j=1}^d |\beta_j|$

$= (\mathbf{y} - \mathbf{X} \boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X} \boldsymbol{\beta}) + \lambda \|\boldsymbol{\beta}\|_1$
Lasso has no closed form.
Bayesian Linear Regression
Setting: Define a prior over the $\boldsymbol{\beta}$ s.
e.g. Ridge:
Assume $\boldsymbol{\beta}$ s distributed with mean 0
 $p(\boldsymbol{\beta}|\Lambda) = \mathcal{N}(\boldsymbol{\beta}|\mathbf{0}, \Lambda^{-1}) \propto \exp(-\frac{1}{2} \boldsymbol{\beta}^T \Lambda \boldsymbol{\beta})$
e.g. Linear Regression:
equivalent to ridge with $\Lambda = \lambda \mathbf{I}, \sigma = 1$
Posterior
given observed \mathbf{X}, \mathbf{y} , use Baye's theorem to find the posterior
 $p(\boldsymbol{\beta}|\mathbf{X}, \mathbf{y}, \Lambda, \sigma) = \mathcal{N}(\boldsymbol{\mu}_{\boldsymbol{\beta}}, \Sigma_{\boldsymbol{\beta}})$
 $\boldsymbol{\mu}_{\boldsymbol{\beta}} = \sigma^2 (\mathbf{X}^T \mathbf{X} + \sigma^2 \Lambda)^{-1} (\mathbf{X}^T \mathbf{y} + \Sigma_{\boldsymbol{\beta}})$
 $\sigma^2 (\mathbf{X}^T \mathbf{X} + \sigma^2 \Lambda)^{-1}$
Bayesian Information Criterion (BIC)
 $-2 \log(\hat{p}(X|\hat{\theta}_k, M_k)) + k' \log n$ tendency to underfit
Akaike Information Criterion (AIC)
 $-2 \log(\hat{p}(X|\hat{\theta}_k)) + 2k', k' = \dim(\theta)$ tendency to select large models (overfit)
Takeuchi Information Criterion (TIC)
 $-2 \log(\hat{p}(X|\hat{\theta}_k)) + 2 \text{trace}[I_1(\theta_k) J_1^{-1}(\theta_k)]$ reduced to AIC if the true model is an element of the model class.
Nonlinear Regression
Idea: Feature space transformation
Model: $\mathbf{Y} = f(\mathbf{X}) = \sum_{m=1}^M \beta_m h_m(\mathbf{X})$
Transformation $h_m(\mathbf{X}) : \mathbb{R}^d \rightarrow \mathbb{R}$
Cubic Spline
e.g. for $d=1$ with knots at ξ_1 and ξ_2
 $h_1(X)=1$ $h_3(X)=X^2$ $h_5(X)=(X-\xi_1)_+^3$
 $h_2(X)=X$ $h_4(X)=X^3$ $h_6(X)=(X-\xi_2)_+^3$
Wavelets
Functions that measure local properties of the underlying data. Keep the most important ones and get rid of the noise.
Gaussian Process Regression
joint Gaussian over all outputs
 $\mathbf{y} = \mathbf{X} \boldsymbol{\beta} + \epsilon$ $\epsilon \sim \mathcal{N}(\epsilon|0, \sigma \mathbf{I}_n)$
We can rewrite the distribution
 $P\left(\begin{bmatrix} \mathbf{y} \\ \mathbf{y}_* \end{bmatrix}\right) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \begin{bmatrix} \mathbf{C}_{\mathbf{y}} & \mathbf{k} \\ \mathbf{k}^T & c \end{bmatrix})$
Such that for prediction:
 $p(\mathbf{y}_*|\mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \mathcal{N}(\mathbf{y}_*|\boldsymbol{\mu}_*, \sigma_*^2)$
 $\boldsymbol{\mu}_{\mathbf{y}_*} = \mathbf{k}^T \mathbf{C}_{\mathbf{y}}^{-1} \mathbf{y}$ $\mathbf{C}_{\mathbf{y}} = \mathbf{K} + \sigma^2 \mathbf{I}$
 $\sigma_*^2 = c - \mathbf{k}^T \mathbf{C}_{\mathbf{y}}^{-1} \mathbf{k}$ $c = k(x_*, x_*) + \sigma^2$
 $\mathbf{k} = k(x_*, \mathbf{X})$
 \mathbf{k} is the kernel function.
lengthscale in kernel: how far can we reliably extrapolate

Bias-Variance tradeoff
 $\text{Bias}(\hat{f}) = \mathbb{E}[\hat{f}] - f^*$
 $\text{Var}(\hat{f}) = \mathbb{E}[(\hat{f} - \mathbb{E}[\hat{f}])^2]$
 $\begin{matrix} |\mathcal{Z}| \downarrow & |\mathcal{F}| \uparrow & \Rightarrow & \text{Var} \uparrow & \text{Bias} \downarrow \\ |\mathcal{Z}| \uparrow & |\mathcal{F}| \downarrow & \Rightarrow & \text{Var} \downarrow & \text{Bias} \uparrow \end{matrix}$
Squared Error Decomposition
 $\mathbb{E}_D \mathbb{E}_{X,Y}[(\hat{f}(X) - Y)^2] = \mathbb{E}_{X,Y}[(\mathbb{E}_Y[Y|X] - Y)^2] \text{ (noise)} + \mathbb{E}_X \mathbb{E}_D[(\hat{f}_D(X) - \mathbb{E}_D[\hat{f}(X)])^2] \text{ (var.)} + \mathbb{E}_X[(\mathbb{E}_D[\hat{f}_D(X)] - \mathbb{E}_Y[Y|X])^2] \text{ (bias}^2)$
(can be derivated by vanishing of the crossproducts)
Parametric Density Estimation
Find the most likely parameter of a distribution.
Maximum Likelihood
Likelihood: $P(\mathcal{X}|\theta) = \prod_{i \leq n} p(x_i|\theta)$
Find: $\hat{\theta} \in \arg \max_{\theta} P(\mathcal{X}|\theta)$
Procedure: solve $\nabla_{\theta} \log P(\mathcal{X}|\theta) = 0$
Efficient & easy to calculate.
Consistent. Converge to best model
 θ_0 Warning: Overfitting!
Maximum A Posteriori
Assume Knowledge of a prior $P(\theta)$
Find: $\hat{\theta} \in \arg \max_{\theta} P(\theta|\mathcal{X}) = \arg \max_{\theta} P(\mathcal{X}|\theta)P(\theta)$
Solve $\nabla_{\theta} \log P(\mathcal{X}|\theta)P(\theta) = 0$
Bayesian Learning
Prior Knowledge of $p(\theta)$
Find Posterior Density: $p(\theta|\mathcal{X})$
Can be done using Baye's Rules
We can use this Recursively:
 $\mathcal{X}^n = \{x_1, \dots, x_n\}$
 $p(\theta|\mathcal{X}^n) = \frac{p(x_n|\theta)p(\theta|\mathcal{X}^{n-1})}{\int p(x_n|\theta)p(\theta|\mathcal{X}^{n-1})d\theta}$ with
 $p(\theta|\mathcal{X}^0)p(\theta)$
Difficult & needs prior knowledge but better against overfitting.
Numerical Est. Techniques
Setting: Estimate $\hat{f}(x) \in \mathcal{F}$ with minimal prediction error.
K-Fold Cross Validation
Initialisation (split training set):
 $\mathcal{Z} = \mathcal{Z}_1 \cup \mathcal{Z}_2 \cup \dots \cup \mathcal{Z}_K, \mathcal{Z}_\mu \cap \mathcal{Z}_\nu = \emptyset$
with map $\kappa : \{1, \dots, n\} \rightarrow \{1, \dots, K\}$
 $|\mathcal{Z}_k| \approx n \frac{K-1}{K}$
Learning:
 $\hat{f}^{-\nu}(x) = \arg \min_{f \in \mathcal{F}} \frac{\sum_{i \in \mathcal{Z}_\nu} (y_i - f(x_i))^2}{|\mathcal{Z} - \mathcal{Z}_\nu|}$
Validation:
 $\hat{R}^{cv} = \frac{1}{n} \sum_{i \leq n} (y_i - \hat{f}^{-\kappa(i)}(x_i))^2$
tendence to Underfit

Leave-one-out: $K = n$ (unbiased but Var can be large \leftarrow corr. datasets)
Bootstrapping
Bootstrap samples: $\mathcal{Z}^* = \{\mathcal{Z}_1^*, \dots, \mathcal{Z}_n^*\}$
each data point in \mathcal{Z}_i^* was randomly drawn from \mathcal{Z} with replacement.
 e_0 Estimator: the error rate for the test data (data that wasn't selected by the bootstrap) is assumed to be the error estimate (e.g. for classification):
 $\hat{R}(S(\mathcal{Z})) = \frac{1}{B} \sum_{b=1}^B \sum_{z_i \notin \mathcal{Z}^*{}^b} \frac{\mathbb{I}_{\mathcal{L}(x_i) \neq y_i}}{|n - \mathcal{Z}^*{}^b|}$
Jackknife
Estimate of an Estimator \hat{S}_n 's Bias.
 $\hat{S}^{JK} = \hat{S}_n - \text{bias}^{JK}$ is JK Estimator.
 $\text{bias}^{JK} = (n-1)(\hat{S}_n - \hat{S}_m)$
 $\tilde{S}_n = \frac{1}{n} \sum_{i=1}^n \hat{S}_{n-1}^{(-i)}$ avg. LOO Estimator.
Debiased est. can have big variance!
Bootstrap **Debiased**
 $\bar{S} = 2\hat{S} - \frac{1}{B} \sum_b \hat{S}^*(b)$
Classification
group points in classes $1, \dots, k, \mathcal{D}, \mathcal{O}$
 \mathcal{D} : doubt class, \mathcal{O} : outliers.
Data: $\mathcal{Z} = \{z_i = (x_i, y_i) : 1 \leq i \leq n\}$ Assume we know $p_y(x) = P[X=x|Y=y]$
Found: classifier $\hat{c} : \mathcal{X} \rightarrow \mathcal{Y} : \{1, \dots, \mathcal{D}\}$
Error: $\hat{R}(\hat{c}|\mathcal{Z}) = \sum_{(x_i, y_i) \in \mathcal{Z}} \mathbb{I}_{\{\hat{c}(x_i) \neq y_i\}}$
Expected Error:
 $\mathcal{R}(\hat{c}) = \sum_{y \leq k} P[y] \mathbb{E}_{x|y}[\mathbb{I}_{\{\hat{c}(x_i) \neq y_i\}} | Y=y]$ (add term from \mathcal{D})
Loss Functions
0-1 Loss: $L^{0-1}(y, z) = \begin{cases} 0 & \text{if}(z=y) \\ 1 & \text{if}(z \neq y) \end{cases}$
Exponential Loss:
 $L^{\text{exp}}(y, z) = \exp(-(2y-1)(2z-1))$
Logistic Loss:
 $L^{\log}(y, z) = \ln(1 + \exp((2y-1)(2z-1)))$
Hinge Loss:
Favors sparsity. Used in SVM
 $L^{\text{hinge}}(y, z) = \max\{0, 1 - (2y-1)(2z-1)\}$
Bayes Optimal Classifier
Minimizes total risk for 0-1 Loss
 $\hat{c}(x) = \begin{cases} y & \text{if } p(y|x) = \max_{z \leq k} p(z|x) > 1-d \\ \mathcal{D} & \text{if } p(y|x) < 1-d \forall y \end{cases}$
Generalize to other loss functions
Discriminant Functions
Functions $g_k(x)$ $1 \leq k \leq K$
Decide: $g_y(x) > g_z(x) \forall z \neq y \Rightarrow$ chose y
Const factor doesn't change decision.
 $g_k(x) = P[y|x] \propto P[x|y]P[y] \Rightarrow g_k(x) = \ln P[x|y] + \ln P[y] = \ln P[x|y] + \pi_y$

Bayes classifier.

Decision Surface of Discriminant

Solve: $g_{k_1}(x) - g_{k_2}(x) = 0$ Special case with Gaussian classes:

if $\Sigma_y = \Sigma \Rightarrow$ linear decision surface

$$g_k(x) = w^T(x - x_0) \quad w = \Sigma^{-1}(\mu_1 - \mu_2)$$

$$x_0 = \frac{1}{2}(\mu_1 + \mu_2) - \frac{\sigma^2(\mu_1 - \mu_2)}{(\mu_1 - \mu_2)^T \Sigma^{-1}(\mu_1 - \mu_2)} \log \frac{\pi_1}{\pi_2}$$

Linear Classifier

optimal for Gaussian with equal cov.

Stat. simplicity & comput. efficiency.

$$g(x) = a^T \tilde{x} \quad a = (w_0, w)^T, \tilde{x} = (1, x)^T$$

$$a^T \tilde{x}_i > 0 \Rightarrow y_i = 1 \quad a^T \tilde{x}_i < 0 \Rightarrow y_i = 2$$

Normalization: $\tilde{x}_i \rightarrow -\tilde{x}_i$ if $y_i = 2$

Find a : $a^T \tilde{x} > 0$ (linearly separable)

Learning w. Gradient Descent:

$$a(k+1) = a(k) - \eta(k) \nabla J(a(k))$$

$J(\cdot)$: cost function $\eta(\cdot)$: learning rate

Newton's rule (opt. grad descent):

$$a(k+1) = a(k) - H^{-1} \nabla J \quad H = \frac{\partial^2 J}{\partial a_i \partial a_j}$$

Perceptron Criterion

$$J_P(a) = \sum_{\tilde{x} \in \tilde{\mathcal{X}}} (-a^T \tilde{x})$$

$\tilde{\mathcal{X}}$ set of misclassified samples.

$\Rightarrow a(k+1) = a(k) + \eta(k) \sum_{\tilde{x} \in \tilde{\mathcal{X}}} \tilde{x}$ Converges if data separable.

WINNOW Algorithm

Performs better when many dimensions are irrelevant. Search for 2 weight vectors a^+, a^- (for each class).

If a point is misclassified:

$$a_i^+ \leftarrow a^{+\tilde{x}_i} a_i^+, a_i^- \leftarrow a^{-\tilde{x}_i} a_i^- \quad (\text{class 1 err.})$$

$$a_i^+ \leftarrow a^{-\tilde{x}_i} a_i^+, a_i^- \leftarrow a^{+\tilde{x}_i} a_i^- \quad (\text{class 2 err.})$$

Exponential update.

Fisher's Linear Discr. Analysis

Maximize distance of the means of the projected classes to find projection plane separating them best.

proj mean: $\tilde{\mu}_\alpha = \frac{1}{n_\alpha} \sum_{x \in \mathcal{X}_\alpha} w^T x = w^T \mu_\alpha$

Dist of proj means: $|w^T(\mu_1 - \mu_2)|$ Class

ses proj. cov: $\tilde{\Sigma}_1 + \tilde{\Sigma}_2 = w^T(\Sigma_1 + \Sigma_2)w$

Fishers Criterion:

$$J(w) = \frac{\|w\|_{\tilde{\Sigma}_1 + \tilde{\Sigma}_2}^2}{\|w\|_{\tilde{\Sigma}_1 + \tilde{\Sigma}_2}^2} = \frac{w^T(\mu_1 - \mu_2)(\mu_1 - \mu_2)^T w}{w^T(\Sigma_1 + \Sigma_2)w}$$

Fishers Crit for Multiple Classes:

$$J(W) = \frac{|W^T S_B W|}{W^T S_W W}$$

$$S_B = \sum_{k=1}^K n_k(\mu_k - \mu)(\mu_k - \mu)^T$$

$$S_W = \sum_{k=1}^K \sum_{x \in \mathcal{D}_i} (x - \mu_i)(x - \mu_i)^T$$

Linear Discriminant for Multiclass

Reformulate as $(k-1)$ "class α - not class α " dichotomie. But some area

are ambiguous

Support Vector Machine (SVM)

Generalize Perceptron with margin and kernel. Find plane that maximizes margin m s.t.

$$z_i g(y) = z_i (w^T y + w_0) \geq m \quad \forall y_i \in \mathcal{Y}$$

$$z_i \in \{-1, +1\} \quad y_i = \phi(x_i)$$

Vectors y_i are the support vectors

Functional Margin Problem:

minimizes $\|w\|$ for $m=1$: $L(w, w_0, \alpha) =$

$$= \frac{1}{2} w^T w - \sum_{i=1}^n \alpha_i [z_i (w^T y_i + w_0) - 1]$$

where α s are Lagrange multipliers.

$\frac{\partial L}{\partial w} = 0$ and $\frac{\partial L}{\partial w_0} = 0$ give us constraints

$$w = \sum_{i=1}^n \alpha_i z_i y_i \quad 0 = \sum_{i=1}^n \alpha_i z_i$$

Replacing these in $L(w, w_0, \alpha)$ we get

$$\tilde{L}(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j z_i z_j y_i^T y_j$$

with $\alpha_i \geq 0$ and $\sum_{i=1}^n \alpha_i z_i = 0$

This is the dual representation. The

optimal hyperplane is given by

$$w^* = \sum_{i=1}^n \alpha_i^* z_i y_i$$

$$w_0^* = -\frac{1}{2} (\min_{z_i=1} w^{*T} y_i + \max_{z_i=-1} w^{*T} y_i)$$

where α maximize the dual problem.

Only Support Vectors ($\alpha_i \neq 0$) contribute to the evaluation.

Optimal Margin: $w^T w = \sum_{i \in \mathcal{SV}} \alpha_i^*$

Discrim.: $g^*(x) = \sum_{i \in \mathcal{SV}} z_i \alpha_i y_i^T y_i + w_0^*$

class = $\text{sign}(y^T w^* + w_0^*)$

Soft Margin SVM

Introduce slack to relax constraints

$$z_i (w^T y_i + w_0) \geq m(1 - \xi_i)$$

$$L(w, w_0, \xi, \alpha, \beta) = \frac{1}{2} w^T w + C \sum_{i=1}^n \xi_i -$$

$$- \sum_{i=1}^n \alpha_i [z_i (w^T y_i + w_0) - 1 + \xi_i]$$

$$- \sum_{i=1}^n \beta_i \xi_i$$

C controls margin maximization vs.

constraint violation

Dual Problem same than usual SVM

but with suppl. constr.: $\alpha_i \leq C$

Non-Linear SVM

use kernel in discriminant funct:

$$g(x) = \sum_{i=1}^n \alpha_i z_i K(x_i, x)$$

E.g solve the XOR Problem with:

$$K(x, y) = (1 + x_1 y_1 + x_2 y_2)^2$$

Multiclass SVM

\forall class $z \in \{1, 2, \dots, M\}$ we introduce

w_z and define the margin m s.t.:

$$(w_z^T y_i + w_{z,0}) - \max_{z \neq z_i} (w_z^T y_i + w_{z,0}) \geq 0$$

$$\forall y_i \in \mathcal{Y}$$

Structured SVM

Each sample y is assigned to a structured output label z

Output Space Representation:

joint feature map: $\psi(z, y)$

scoring function: $f_w(z, y) = w^T \psi(z, y)$

$$\hat{z} = h(y) \arg \max_{z \in \mathcal{K}} f_w(z, y)$$

Kernels

Similarity based reasoning

$$\text{Gram Matrix } K = (K(x_i, x_j)) \quad 1 \leq i, j \leq n$$

$$K(x, x') = \phi(x)^T \phi(x') \quad K(x, x') = K(x', x)$$

$K(x, x')$ pos. semi-def. (all EV ≥ 0)

If K_1 & K_2 are kernels K is too:

$$K(x, x') = K_1(x, x') K_2(x, x')$$

$$K(x, x') = \alpha K_1(x, x') + \beta K_2(x, x')$$

$$K(x, x') = K_1(h(x), h(x')) \quad h: \mathcal{X} \rightarrow \mathcal{X}'$$

$$K(x, x') = h(K_1(x, x')) \quad h: \text{poly/exp}$$

Kernel Function Examples:

$$K(x, x') = x^T x' \quad K(x, x') = (x^T x' + 1)^p$$

$$\text{RBF (Gauss): } K(x, x') = \exp(-\|x - x'\|^2 / 2h^2)$$

$$\text{Sigmoid: } K(x, x') = \tanh(\alpha x^T x' + c)$$

not p.s.d eg: $x = [1, -1], x' = [-1, 2]$

Ensemble Methods

Combining Regressors

set of estimators: $\hat{f}_1(x), \dots, \hat{f}_B(x)$ simple

$$\text{average: } \hat{f}(x) = \frac{1}{B} \sum_{i=1}^B \hat{f}_i(x)$$

$$\text{Bias}[\hat{f}(x)] = \frac{1}{B} \sum_{i=1}^B \text{Bias}[f_i(x)]$$

$$\text{Var}[\hat{f}(x)] \approx \frac{\sigma}{B} \text{ if the estimators are uncorrelated.}$$

Combining Classifiers

Input: classifiers $c_1(x), \dots, c_B(x)$

$$\text{Infer } \hat{c}_B(x) = \text{sgn}(\sum_{b=1}^B \alpha_b c_b(x))$$

with weights $\{\alpha_b\}_{b=1}^B$

Requires diversity of the classifiers.

Bagging

Train on bootstrapped subsets.

Sample: $\mathcal{Z} = \{(x_1, y_1), \dots, (x_n, y_n)\}$

\mathcal{Z}^* : chose i.i.d from \mathcal{Z} w. replacement

Random Forest (Bagging strategy)

Collection of uncorr. decision trees.

Partition data space recursively.

Grow the tree sufficiently deep to reduce bias.

Prediction with voting.

Boosting

Combine uncorr. weak learners in sequence. (Weak to avoid overfitting).

Coeff. of \hat{c}_{b+1} depend on \hat{c}_b 's results

AdaBoost (minimizes exp. loss)

Init: $\mathcal{X} = \{(x_1, y_1), \dots, (x_n, y_n)\}, w_i^{(1)} = \frac{1}{n}$

Fit $\hat{c}_b(x)$ to \mathcal{X} weighted by $w^{(b)}$

$$\epsilon_b = \sum_{i=1}^n w_i^{(b)} \mathbb{I}_{\{c_b(x_i) \neq y_i\}} / \sum_{i=1}^n w_i^{(b)}$$

$$\alpha_b = \log \frac{1 - \epsilon_b}{\epsilon_b}$$

$$w_i^{(b+1)} = w_i^{(b)} \exp(\alpha_i \mathbb{I}_{\{c_b(x_i) \neq y_i\}})$$

return $\hat{c}_B(x) = \text{sgn}(\sum_{b=1}^B \alpha_b c_b(x))$

best approx. at log-odds ratio.

Neural Networks

Multi Layer Perceptron

$\{x_j\}_{j=1}^J$ input, $\{y_i\}_{i=1}^I$ output

$\{z_k^l\}_{k=1}^{K(l)}$ hidden nodes in layer l $1 \leq l \leq L$

$w_m k^l$ weights from z_k^{l-1} to z_m^l

$w_i k^{L+1}$ weights from z_k^L to output y_i

$$z_k^l = h(a_k^l) = h(\sum_{m=1}^{K(l-1)} w_{km}^l z_m^{l-1})$$

$$y_i = \sigma(a_i^{L+1}) = h(\sum_{m=1}^{K(L)} w_{im}^{L+1} z_m^L)$$

$$\mathcal{L}(\hat{y}(\mathbf{W}, \mathbf{X}), y) = \sum_{n=1}^N \mathcal{L}_n(\hat{y}(\mathbf{W}, \mathbf{X}_n), Y_n)$$

$L = 0$ or $h(a) = a \Rightarrow$ multiple lin. reg.

Layers \Rightarrow generaliz. & simplicity.

Model data generating mechanism.

Backpropagation

Effc. evaluation of loss derivative:

$$\frac{\partial \mathcal{L}_n}{\partial w_{ik}^{L+1}} = \delta_i^{L+1} z_k^L \quad \frac{\partial \mathcal{L}_n}{\partial w_{mk}^l} = \delta_m^l z_k^{l-1}$$

$$\delta_i^{L+1} = (\hat{y}_i - y_i) \sigma'(\sum_{m=1}^{K(L)} w_{im}^{L+1} z_m^L)$$

$$\delta_m^l = (\sum_{r=1}^{K(l+1)} \delta_r^{l+1} w_{rm}^{l+1})$$

$$h'(\sum_{r=1}^{K(l-1)} w_{mr}^{l-1} z_r^{l-1})$$

$$w_{ij}^l \leftarrow w_{ij}^l + \eta \delta_i^l z_j^{(l-1)}$$

Regularization

Avoid overfitting on complex nets.

Early Stopping separate data into

train/error/validation sets.

Drop Out Combine thinned nets

with removed nodes.

Bayesian priors on w 's

Autoencoder

Data compression purposes, Output

should reproduce input. \Rightarrow PCA

Convolutional Neural Network

Modelling invariance. Convolutional

Layers (filters on a region) & Pooling

Layers (aggregate nodes together).

Unsupervised Learning

Histograms

$$p_i = \frac{n_i}{N \Delta_i} \quad n \leq N \text{ in bin } i \text{ of size } \Delta_i$$

Not scaling to multiple dimensions.

$$K \simeq NP \quad P \simeq p(x) V \Rightarrow p(x) = \frac{K}{N V}$$

K #samples in region of volume V , P

probability of falling in it.

Kernel Density Estimator

Fix V and determine K .

$$\text{Gaussian Kernel: } \phi(u) = \frac{\exp(-\frac{1}{2}\|u\|^2)}{\sqrt{2\pi}}$$

Result in a smoother density model

$$p(x) = \frac{1}{N} \sum_{n=1}^N \frac{1}{(2\pi h^2)^{D/2}} \exp(-\frac{\|x - x_n\|^2}{2h^2})$$

We can chose any other kernel ϕ with

$$\phi(u) \geq 0 \quad \int \phi(u) du = 1$$

K-Nearest Neighbors

Fix K and find V

$$\hat{p}(x) = \frac{1}{V_k(x)}, v_k(x) \text{ minimal volume}$$

around x containing k neighbors.

Classifier: classify x by the majority

of the vote of its k -NN.

1-NN Error Rate the 1-NN error rate

P is always $P^* \leq P \leq 2P^*$ where P^* is

the error rate of the Bayes rule. \Rightarrow as

k goes to infinity kNN becomes optimal

kNN not optimal if class densities

are very different.

Mixture Models

Gaussian Mixture

EM-Algorithm

Latent Variable: unknown data \rightarrow

What cluster generated each sample?

EM does ML for unknown parameters.

Latent var. $M_{xc} = \begin{cases} 1 & \text{c generated x} \\ 0 & \text{else} \end{cases}$

$$P(\mathcal{X}, M | \theta) = \prod_{x \in \mathcal{X}} \prod_{c=1}^k (\pi_c P(x | \theta_c))^{M_{xc}}$$

E-Step

$$\gamma_{xc} = \mathbb{E}[M_{xc} | \mathcal{X}, \theta^{(j)}] = \frac{P(x | c, \theta^{(j)}) P(c | \theta^{(j)})}{P(x | \theta^{(j)})}$$

M-Step

$$\mu_c^{(j+1)} = \frac{\sum_{x \in \mathcal{X}} \gamma_{xc} x}{\sum_{x \in \mathcal{X}} \gamma_{xc}}$$

$$(\sigma_c^2)^{(j+1)} = \frac{\sum_{x \in \mathcal{X}} \gamma_{xc} (x - \mu_c)^2}{\sum_{x \in \mathcal{X}} \gamma_{xc}}$$

$$\pi_c^{(j+1)} = \frac{1}{|\mathcal{X}|} \sum_{x \in \mathcal{X}} \gamma_{xc}$$

k-Means

identify clusters of data.

Given $\mathcal{X} = \{x_1, \dots, x_n\}$

Find $c(\cdot)$ and \mathcal{Y} minimizing

$$\mathcal{R}^k m(c, \mathcal{Y}) = \sum_{x \in \mathcal{X}} \|x - \mu_{c(x)}\|^2$$

Assign to nearest cluster. Recompute all

clusters and repeat. Also called hard

EM. Special case of GMM w. uniform

prior and diag. covariance ($\rightarrow 0$).

Extras

Taylorreihe: $\sum_{n=0}^{\infty} \frac{f^{(n)}(a)}{n!} (x - a)^n$

Convex/Concav: $f' \geq 0$ or $f' \leq 0$

LinAlg: $X_{-i} Y_{-j}^T = X Y^T - x_i y_i^T$