

# Parallele Programmierung ZF

Equivalent if  $H1A == G1A$  and  $H1B == G1B$

## O-Notation

$O(\pm) = (\text{work} / \text{span})$

Work = work done  $O(\text{linear})$

span =  $\sum_{i=1}^n 1/p + \text{inf} = O(\text{lin.})/p + O(\text{inf})$

## Parallelism

Distributed Mem.  $\rightarrow$  no problem

Shared Mem.  $\rightarrow$  data races etc.

- Simultaneous Multithreading SMT
- Multicores
- Symmetric Multicore Processing SMP
- Non Uniform Memory Access NUMA

## Distributed Data

+ No coordination

- Messages for data access

## Shared Data

- Mutual exclusion (Mutex)

## Task Parallel

- Programmer defines Tasks
- Generic

## Data Parallel

- Set of Operations is applied simultaneously on individual items.
- Work partitioning by System
- $\rightarrow$  programmer decides what, not how

## Bad Interleaving

- "high level race condition"
- Problem with unfavorable execution

## Data Races

- "low level race condition"
- If two threads access the same memory location  $\rightarrow$  data inconsistency
- Not every instruction is executed in one step

avoid this: use one of the following three

- thread local memory
- Immutable memory  $\rightarrow$  do not write/change
- Synchronized

## Correctness

Mutual Exclusion

Freedom from Deadlock

Freedom from Starvation

Lock-free // Wait-free

	non Blocking	Blocking
all make progress	Wait-free	Starvation-free
1 makes progress	lock-free	Deadlock-free

## Deadlock

Two or more processes are mutually blocked because each process waits for an other to proceed

## Starvation

Repeated but unsuccessful attempt of recently unblocked process to continue its execution

## Livelock

Competing processes are able to detect a deadlock but make no progress solving it.

## Lock-free

At least one makes progress ( $\Rightarrow$  not starvation free)

## Wait-free

All make progress  $\Rightarrow$  Starvation free

## avoid Deadlocks

create unique order, get one with smaller first.

## Single Write Multiple read

## Validation

- Find Elements, • load them,
- Check condition again, • If it holds, done else repeat

## Lazy list Remove

- Find nodes • lock them
- Mark the one to be deleted • redirect pointer

## Lock-free-programming

- read data • modify data
- Compare and Set • if CAS returns false repeat

## Producer and Consumer

- A producer does sth with the data and enqueues it to the array.
- A consumer takes it from there and continues working with it.

## Datastructure

Circular array with enqueue and dequeue pointer  
in = ... out = ...

## Coordination

barrier, send/receive

## Linearizability

$(\rightarrow G) \subseteq (\rightarrow S)$  is required

If it is possible to arrange the events on the timeline, so that a given output is true

Invocation A:  $g.do(x)$

Result A:  $g.result$

## History

is linearizable if possible to extend it to a legal sequential history

$\rightarrow$  Sequence of invocations and Results

pending: if invocation does not process a result.

Complete Subhistory: if not pending

## Sequential History

- Method calls of different threads do not interleave
- final pending is ok

## Well formed History

- if proj. per thread are the same.

## Legal History

- if proj. on every obj. is sequential
- Methods preceding, overlapping

## Projections

Objects A:  $g.do()$ ;  
A:  $g.result()$ ;  
B:  $g.do()$ ;

Threads B:  $g.do()$ ;  
B:  $g.result()$ ;  
B:  $e.do()$ ;  
B:  $e.result()$ ;

## Realtime Order

$(\rightarrow G) \subseteq (\rightarrow S)$

$(\rightarrow G) = \{a \rightarrow c, b \rightarrow c\}$

$(\rightarrow S) = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$

## Sequential Consistency

If G can be extended to legal sequential History S for each thread

- Timelines können gegeneinander verschoben werden!

$(\rightarrow G) \subseteq (\rightarrow S)$  is not required

Operations need to respect program order but not realtime order

- possible to reorder operations done by different threads

Power-supply-wall CPU can't get enough power from PSU

Power-dissipation-wall CPU can't dissipate heat  $\rightarrow$  overheating

## Transactional Memory (Hazard Pointers)

Summarizes executions to transactions. If two transactions access to the same memory element. One transaction gets reversed.

## Spinlock

- No notification mechanism.
- Computing resources wasted, overall performance degraded, particularly for long-lived contention
- Scheduling fairness / missing FIFO behaviour. Solved with queue locks

## Backoff

sleep or wait instead of while for waiting

## Bakery Algorithm (Ticket-system) FIFO

## Peterson Algorithm

- satisfies mutual exclusion
- starvation free

## Programming Model: What needs to be done (declarative)



**Moore's law** - 2005: transistors  $\times 2$  for 2 years  
 - Exponentially faster  
 Problem: Heat + Power

### Single-Core

- + Flexibility
- + Performance
- Complex Control HW
- Expensive Power

### Multi-Core

- + Simpler Control HW
- + More data throughput
- + More Power efficient
- complex programming

Hardware

**Key Questions:** always? why? better?  
 sort by quality? Termination?  
 instructions unambiguous?

### Prove correctness

- Induction
- Invariant // assertion

### JVM Java Virtual Machine

- Interprets compilation code
- Simulates CPU
- Translates compilation code to machine code

### Performance loss with II

- Context switches
- loss of locality
- CPU scheduling time vs. running time
- Additional overhead with synchronisation

**Concurrent:** requests for <sup>limited</sup> resources, manage access to "synchronisation" shared resources (possible on <sup>single cycle</sup> multi-core)

**Parallelism:** work on resources, use extra resources to "coordinate" solve a problem faster

**Principle of locality** easier to access local data

### Sorts of Parallelism

- Vectorisation  $n$  times same in parallel
- Pipelining same thing but at different states
- ILP: Instruction level Parallelism
  - Pipelining
  - Super Scalar CPU (mult. Instr. per cycle)
  - Out of order execution
  - Speculative execution

**Throughput:** Input // Output data rate (1 / längste Zeiteinheit)

**Latency:** Time for one object to pass

1 Obj.: Add all steps

$n$  Obj.: Overall time / # Obj.

**Balanced Pipelines** every unit  $\rightarrow$  same timestep

### Execution Time:

$T_1$  = sequential Time

$T_p \geq T_1 / p$

**Speed-Up:**  $S_p = T_1 / T_p \leq p$

**Efficiency:**  $S_p / p$

fixed amount of work

**Amdahl's law** problem size is constant

$$T_1 = W_{ser} + W_{par}$$

$$T_p \geq W_{ser} + W_{par} / p, T_p \geq \text{crit. path}$$

$$W_{ser} = f \cdot T_1, W_{par} = (1-f) T_1$$

$$S_p \leq \frac{1}{f + \frac{1-f}{p}}, f = \text{non parallelizable serial fractions}$$

- + maximum speedup  $a_p = \text{Overhead}$
- = limit on scalability

- all non-parallel parts can cause problems!

**Gustafson's law** runtime is constant  $T_1 = T_{inf}$

$$W = p(1-f)T_1 + fT_1$$

$$S_p = f + p(1-f) = p - f(p-1)$$

**Reentrant lock** (acquire lock if already owns)

lock.acquire

lock.release

... finally { lock.unlock(); }

Condition

Final condition Not Full = lock.newCondition();

NotFull.await();

signal();

signalAll();

while (!isFull) { await(NotFull); lock.unlock(); }

**Synchronized (Expression) { ... }**

synchron... (Obj.) { ... }

synchron... (this) { ... }

synchron... (lock) { ... }

synchron... methodname () { ... }

... methodname () { ... synchron... (int ?) { ... } }

### Consensus

You cannot implement lock-free algorithms that require CAS with atomic registers.

### Message Passing Interface (MPI)

- Hides Software / Hardware details
- Portable, flexible
- Implemented as a library

### Synchronous Message Passing

- sender blocks until message is received

### Asynchronous Message Passing

- sender does not block (fire and forget)
- placed into a buffer for receiver to get

### Bitonic Sort

Parallel algorithm for sorting

Sequential:  $O(n \log^2 n)$

Parallel:  $O(\log^2 n)$

### Read Modify Write

**TAS** Test-And-Set returns boolean

if (mem.ref[s] = 0) {

mem[s] = 1;

return true; }

else return false; eg: while (!TAS(integer))...

**CAS** Compare-And-Swap Java: Boolean, Else: old value

atomic int i

i = compareAndSet(j) eg. while (!CAS(lock, 0, 1))

if (old = oldval) {

old = a;

return true; }

else return false;

boolean attemptMark (V expected ref., boolean newMark)

boolean compareAndSet (V expected ref., V new ref., boolean expMark, boolean newMark)

**TTAS** test Test And Set

public void lock () {

do

while (state.set()) { }

while (!state.compareAndSet(false, true)) { }

public void unlock () { state.set(false); }

**Semaphores** atomic counter (can be acquired if  $> 0$ )

acquire (s) {

wait until  $S > 0$ ;

dec (s); }

release (s) { inc (s); }

### Monitors

allows mutual exclusion and wait() synchronized in JAVA

wait ();

notify ();

notifyAll ();

### Volatile

private volatile int x = 0; ...

x is treated like a synchronized and written into memory directly

### Oblivious

"Starts knowing nothing", makes the same comparisons regardless of the input.

### Redundant

"Überflüssig", expendable comparisons



**Interface** No code sharing

```
public interface Shape {
    public double area();
    public double perimeter();
}

public class Rectangle implements Shape {
    public Rectangle() {}
    @Override
    public double area() {}
    @Override
    public double perimeter() {}
}
```

**Runnable**

```
class
public MyRunnable implements Runnable {
    public MyRunnable...
    @Override
    public void run() {}
}

public class Use {
    Thread t1 = new Thread(new MyRunnable());
    t1.start();
    Thread[] t = new Thread[5];
    t[0] = new Thread(new MyRunnable());
    t1.join();
}
```

**Extend Thread**

```
public class MyThread extends Thread {
    @Override
    public void run() {}
}
```

```
public class Use {
    MyThread t1 = new MyThread();
    t1.start();
    t1.join();
}
```

**Tasks:**

fast context switch, cheap to create, can create many of them to enable finegrained parallelism, matches with short lifetime of web requests

**Threads:**

more expensive to create, reuse a pool of threads (sized to match #cores), apply work stealing to achieve load balancing.

**Callable** interface with return type

```
public class MyCallable implements Callable<Integer> {
    public MyCallable() {}
    @Override
    public Integer call() {}
}

public class Use {
    public static void main(String args[]) {
        Callable<Integer> thisCallable = new MyCallable();
        ExecutorService pool = Executors.newFixedThreadPool();
        Future<Integer> result = pool.submit(thisCallable);
        int a = result.get();
        pool.shutdown();
    }
}
```

**Fork / Join**

```
public class fjTask extends RecursiveTask<V> {
    @Override
    public V compute() {
        if small enough do it yourself
        else
            RecursiveTask<V> l = new fjTask(...);
            " " " r = " "
            l.fork();
            V a = r.compute();
            V b = l.join();
            return V;
    }
}
```

```
public class Use () {}
... {
    ... static ForkJoinPool f1 = new ForkJoinPool();
    Result = f1.invoke(new fjTask(input));
}
```

**Reduce Lock Contention**

- Reduce duration locks are held (smaller synchronized blocks)
- Reduce frequency of lock requests (lock splitting or stripping)
- Replace exclusive locks with coordination mechanisms that permit greater concurrency (R/W locks, non-bl. DS)
- Split up a lock into smaller locks (fine-grained locking)
- Avoid locks entirely (lock-free, replication, immutability)
- Transactional memory
- fairness model

**Barrier** Rendezvous for arbitrary number of threads  
init(N), await(), reset()

eg.: Game of life, displaying the pixels in an image.

**Threads +1** (share the same address)

Terminates only if run() terminates

```
t1.start(); // creates a new Thread
    .run(); // same Thread do run()
    .join(); // = finish
Backoff {
    .sleep(); // use in while blocks
    .notifyAll();
    .setPriority(); // 1-10
    .setName("...");
    .getId();
    .getState(); // new, runnable, blocked, waiting, terminated
    .MAX_PRIORITY // MIN-P...
    .interrupt(); // can be ignored
    .isInterrupted(); // now?
    .interrupted(); // Present
```

**Threadpool**

↳ Callable

use seq. cutoff 500-1000

**Concurrent**

- + Concurrent collection allows multiple threads to be reading / modifying it in parallel
- hard to make compound operations atomic, since you cannot use client-side locking

**Problem:** Load balancing **Solution:** Dynamic Scheduling

**API design properties**

declarative, no storage, functional, <sup>easy</sup> to parallelize  
less flexibility than for-loop, operators must be stateless.

**Data-Parallelism**

+ maps, reductions, filter

- complex synchronizations, computation with complex computation patterns

**Multi-threaded Programs**

Are hard to test because of nondeterminism from scheduling. Hard to track down bugs (subsequent runs)

**Problems in programs using locks**

- performance, lock contention
- Deadlocks
- Priority Inversion
- Livelock



```

class ReadersWritersLock {
    int writers = 0;
    int readers = 0;
    int writersWaiting = 0;

    • synchronized void AcquireRead () {
        while (writers > 0 || writersWaiting > 0) {
            try { wait (); }
            catch (InterruptedException e) {} }
        readers ++; }

    • synchronized void ReleaseRead () {
        readers --;
        notifyAll (); }

    • synchronized void AcquireWrite () {
        writersWaiting ++;
        while (writers > 0 || readers > 0) {
            try { wait (); }
            catch (InterruptedException e) {} }
        writersWaiting --;
        writers ++; }

    • synchronized void ReleaseWrite () {
        writers --;
        notifyAll (); }
    
```

= set equal  
 == equal?  
 a.equals(b) returns boolean

String a

- a.indexOf (str);
- length ();
- substring (index 1, index 2);
- toLowerCase ();
- toUpperCase ();
- equals (string);
- equalsIgnoreCase (string);
- startsWith (str);
- endsWith (str);
- contains (str);

```

public class NameClass {

    public NameClass (int i, ... ) {} {
        // Define "Creation Method"
    }

    NameClass NameObject =
        new NameClass (parameters);
}

public class Use {
    int a, b, c, x1, flavour;

    public static void main (String args []) {
        [ if (...) { ... }
        [ if else (...) { ... }
        [ else { ... }
        - System.out.println (a + "..." + "/N");
        - while (...) { }
        [ do { ... }
        [ while { ... }
        Scanner s = new Scanner (System.in);
        int d = s.nextInt();
        String e = s.nextLine();
        e = s.lastString (); }

    public static int f (int i) {
        return i; }
    }
    
```

```

public class dog extends animal() {
    redefine ...
    @Override
    ... }

    for (int i = 0; i < 100; i++) { }

    switch (Data) {
        case 1: ...
        case 2: ... }
    
```

```

import static java.lang.Math.*;

Math.min (last, first)

System.out.println (...);
    
```